

NovelPerspective

Anonymous ACL submission

Abstract

We present NovelPerspective: a proof of concept tool which allows consumers to subset ebook novels, based on the main character of the section. Many novels have multiple main characters each with their own storyline running in parallel. A well-known example is George R. R. Martin’s novel: “A Game of Thrones”, and the others from that series. Our tool detects which character each section is about, and allows the user to generate a new ebook with only those sections. This gives consumers new options in how they consume their media; allowing them to pursue the storylines sequentially, or skip chapters about characters they find boring. We present two simple baselines, and several machine learning based methods for the detection of the main character.

1 Introduction

Often each section of a novel is written from the perspective of a different main character. The characters each take turns in the spot-light, with their own parallel storylines being unfolded by the author. Personally, as readers, we have often desired to read just one storyline at a time, particularly when reading the book a second-time. We present a tool, NovelPerspective, to give the consumer this choice.

Our tool allows the consumer to select which characters of the book they are interested in, and to generate a new ebook file containing just the sections from that character’s point of view (POV). The critical part of this system is the detection of the POV character. This is not an insurmountable task, building upon the well established field of named entity recognition. However to our knowl-

edge there is no software to do this. Such a tool would have been useless, in decades past when books were distributed only on paper. But today, the surge in popularity of ebooks has opened a new niche for consumer narrative processing. Tools such as the one presented here, give the reader new freedoms in controlling how they consume their media.

Having a large cast of characters, in particular POV characters, is a hallmark of the epic fantasy genre. Well known examples include: George R.R. Martin’s “A Song of Ice and Fire”, Robert Jordan’s “Wheel of Time”, Brandon Sanderson’s “Cosmere” universe, Steven Erikson’s “Malazan Book of the Fallen”, amongst thousands of others. Generally, these books are written in *limited* third-person POV; that is to say the reader has little or no more knowledge of the situation described than the main character does.

We focus here on novels written in the *limited* third-person POV. In these stories, the main character is, for our purposes, the POV character. Limited third-person POV is written in the third-person, that is to say the character is referred to by name, but with the observations limited to being from the perspective of that character. This is in-contrast to the *omniscient* third-person POV, where events are described by an external narrator. Limited third-person POV is extremely popular in modern fiction. It preserves the advantages of first-person, in allowing the reader to observe inside the head of the character, while also allowing the flexibility to the perspective of another character (Booth, 2010). This allows for multiple concurrent storylines around different characters. Our tool helps users un-entwine such storylines, giving the option to process them sequentially.

The utility of dividing a book in this way varies with the book in question. Some books will cease to make sense when the core storyline crosses over

different characters. Other novels, particularly the epic fantasy stories which we are primarily considering, have many parallel storylines which only rarely intersect. While we are unable to find a formal study on this, anecdotally many readers speak of:

- “Skipping the chapters about the boring characters”
- “Only reading the *real* main character’s sections”
- “Reading ahead, past the side-stories, to get on with the main plot”

Particularly if they have read the story before, and thus do not risk confusion. Such opinions are a matter of the consumer’s personal taste. The NovelPerspective tool gives the reader the option to customise the book in this way, according to their personal preference.

We note that sub-setting the novel once does not prevent the reader from going back and reading the intervening chapters if it ceases to make sense, or from sub-setting again to get the chapters for another character whose path intersects with the storyline they are currently reading. We can personally attest for some books reading the chapters one character at a time is indeed possible, and pleasant: the first author of this paper read George R.R. Martin’s “A Song of Ice and Fire” series in exactly this fashion.

The primary difficulty in segmenting ebooks this way attributing each section to its POV character. That is to say detecting who is the point of view character. Very few books indicate this clearly, and the reader is expected to infer it during reading. This is easy for most humans, but automating it is a challenge. To solve this, the core of our tool is its character classification systems. We investigated several options which the main text of this paper will discuss.

2 Character Classification Systems

The full NovelPerspective pipeline is shown in Figure 1. The core character classification step (step 3), is detailed in Figure 2. In this step the raw text is first enriched with parts of speech, and named entity tags. From this, features are extracted for each named entity. These feature vectors are used to score the entities for the most-likely POV character. The highest scoring charac-

ter is returned by the system. The different systems presented modify the **Feature Extraction** and **Character Scoring** steps. A broadly similar idea, for detecting the focus location of news articles, was presented by (Imani et al., 2017).

2.1 Baseline systems

To the best of our knowledge no systems have been developed for this task before. As such, we have developed two deterministic baseline character classifiers. These are both potentially useful to the end-user in our deployed system (Section 5), and used to gauge the performance of the more complicated systems in the evaluations presented in Section 4.

It should be noted that the baseline systems, while not using machine learning for the character classification steps, do make extensive use of machine learning-based systems during the preprocessing stages. The POS-tagger, and the Named Entity recogniser are based on machine learning.

2.1.1 “First Mentioned” Entity

An obvious way to determine the main character of the section is to select the first named entity. We use this to define the “First Mentioned” baseline. In this system, the **Feature Extraction** step is simply retrieving the position of the first use of each name; and the **Character Scoring** step assigns each a score such that earlier is higher. This works well for many examples: “*One dark and stormy night, Bill heard a knock at the door.*”; however it fails for many others “*‘Is that Tom?’ called out Bill, after hearing a knocking.*”. Sometimes a section may go several paragraphs describing events before it even mentions the character who is perceiving them. This is a varying element of style.

2.1.2 “Most Mentioned” Entity

A more robust method to determine the main character, is to use the occurrence counts. We call this the “Most Mentioned” baseline. The **Feature Extraction** step is to count how often the name is used. The **Character Scoring** step assigns each a score based what proportional of all names used were for this entity. This works well for many books. The more important a character is, the more often their name occurs. However, it is fooled, for example, by book chapters that are about the POV character’s relationship with a secondary character. In such cases the secondary character may be mentioned more often.




Figure 1: The full NovelPerspective pipeline. Note that step 5 uses the original ebook to subset.




Figure 2: The general structure of the character classification systems. This is step 3 of the whole pipeline shown in Figure 1.

2.2 Machine learning systems

One can see the determination of the main character as a multi-class classification problem. From the set of all named entities in the section, classify that section as to which one is the main character. Unlike typical multi-class classification problems the set of possible classes varies per section being classified. Further, even the total set of possible named characters, i.e. classes, varies from book to book. An information extraction approach is required which can handle these varying classes. As such, a machine learning model for this task can not incorporate direct knowledge of the classes (i.e. character names).

We reconsider the problem as a series of binary predictions. The task is to predict if a given named entity is the point of view character. For each possible character (i.e. each named-entity that occurs), a feature vector is extracted (see Section 2.2.1). This feature vector is the input to a binary classifier, which determines the probability that it represents the main character. The **Character Scoring** step is thus the running of the binary classifier: the score is the output probability normalised over all the named entities.

2.2.1 Feature Extraction for ML

We investigated two feature sets as inputs for our machine learning-based solution. They correspond to different **Feature Extraction** steps in Figure 2. A hand-engineered feature set, that we call the “Classical” feature set; and a more modern “Word Embedding” feature set. Both feature sets give information about how the each named entity token was used in the text.

The “Classical” feature set uses features that are well established in NLP related tasks. The features can be described as *positional features*, like in the First Mentioned baseline; *occurrence count fea-*

tures, like in the *Most Mentioned* baseline and *adjacent POS counts*, to give usage context. The *positional features* are the index (in the token counts) of the first and last occurrence of the named entity. The *occurrence count features* are simply the number of occurrences of the named entity, supplemented with its rank on that count compared to the others. The *adjacent POS counts* are the counts of each part of speech tag on the word prior to the named entity, and on the word after. We theorised that this POS information would be informative, as it seemed reasonable that the POV character would be described as doing more things, so co-occurring with more verbs. This gives 100 base features. To allow for text length invariance we also provide each of the base features expressed as a portion of its maximum possible value (e.g. for a given POS tag occurring before a named entity, the portion of times this tag occurred). This gives a total of 200 features.

The “Word Embedding” feature set uses FastText word vectors (Bojanowski et al., 2016). We concatenate the word embedding for the word immediately prior to, and immediately after each occurrence of a named entity; and take the element-wise mean of this concatenated vector over all occurrences of the entity. Such averages of word embeddings have been shown to be a useful feature in many tasks (White et al., 2015; Mikolov et al., 2013). This has a total of 600 features.

2.2.2 Classifier

The binary classifier, that predicts if a named entity is the main character, is the key part of the **Character Scoring** step for the machine learning systems. From each text in the training dataset we generated a training example for every named entity that occurred. All but one of these was a negative example. We then trained it as per normal

for a binary classifier. The score for a character is the classifier’s predicted probability of its feature vector being for the main character.

Our approach of using a binary classifier to rate each possible class, may seem similar to the one-vs-rest approach for multi-class classification. However, there is an important difference. Our system only uses a single binary classifier; not one classifier per class, as the classes in our case vary with every item to be classified. The fundamental problem is information extraction, and the classifier is a tool for the scoring which is the correct information to report.

For the Classical Feature-set, we used logistic regression, with the features being preprocessed with 0-1 scaling. During preliminary testing we found that many classifiers had similar high degree of success, and so chose the simplest. With the higher dimensional feature-sets we used a radial bias support vector machine, with standardisation during preprocessing, as has been commonly used with word embeddings on other tasks.

3 Experimental Setup

3.1 Datasets

We make use of three series of books selected from our own personal collections. The first four books of George R. R. Martin’s “A Song of Ice and Fire” series (hereafter referred to as ASOIAF); The two books of Leigh Bardugo’s “Six of Crows” duology (hereafter referred to as SOC); and the first 9 volumes of Robert Jordan’s “Wheel of Time” series (hereafter referred to as WOT). In Section 4 we consider the use of each as a training and testing dataset. In the online demonstration (Section 5), we deploy models trained on the combined total of all the datasets.

To use a book for the training and evaluation of our system, we require a ground truth for each section’s POV character. ASOIAF and SOC provide ground truth for the main character in the chapter names. Every chapter only uses the POV of that named character. WOT’s ground truth comes from an index created by readers¹. We do not have any datasets with labelled sub-chapter sections, though the tool does support such works.

The total counts of chapters and characters in the datasets, after preprocessing, is shown in Table 1. Preprocessing consisted of discarding chap-

¹http://wot.wikia.com/wiki/List_of_Point_of_View_Characters

Dataset	Chapters	POV Characters
ASOIAF	256	15
SOC	91	9
WOT	432	52
combined	779	76

Table 1: The number of chapters and point of view characters for each dataset

ters for which the POV character was not identified (e.g. prologues); and of removing the character names from the chapter titles as required.

3.2 Evaluation Details

In the evaluation, the systems are given the body text and asked to predict the character names. During evaluation, we sum the scores of the characters alternative aliases/nick-names used in the books. For example merging Ned into Eddard in ASOIAF. This roughly corresponds to the case that a normal user can enter multiple aliases into our application when selecting sections to keep. We do not use these aliases during training, though that is an option that could be investigated in a future work.

3.3 Implementation

The full implementation is available at <https://github.com/oxinabox/NovelPerspective/>. The text is preprocessed using NLTK (Bird and Loper, 2004). It is tokenised, POS tagged, and finally named entity chunked, using NLTK’s default methods. Specifically, these are the Punkt sentence tokenizer, the regex-based improved TreeBank word tokenizer, greedy averaged perceptron POS tagger, and the max-entropy binary named entity chunker. The use of a binary, rather than a multi-class, named entity chunker is significant. Fantasy novels often use “exotic” names for characters, we found that this often resulted in character named entities being misclassified as organisations or places. Note that this is particularly disadvantageous to the First Mentioned baseline, as any kind of named entity will steal the place. Nevertheless, it is required to ensure that all character names are a possibility to be selected.

Scikit-Learn is used for the machine learning and evaluations (Pedregosa et al., 2011).

Test Set	Method	Train Set	Acc
ASOIAF	First Mentioned	—	0.250
ASOIAF	Most Mentioned	—	0.914
ASOIAF	ML Classical Features	SOC	0.953
ASOIAF	ML Classical Features	WOT	0.984
ASOIAF	ML Word Emb. Features	SOC	0.863
ASOIAF	ML Word Emb. Features	WOT	0.977
SOC	First Mentioned	—	0.429
SOC	Most Mentioned	—	0.791
SOC	ML Classical Features	ASOIAF	0.923
SOC	ML Classical Features	WOT	0.923
SOC	ML Word Emb. Features	ASOIAF	0.945
SOC	ML Word Emb. Features	WOT	0.934
WOT	First Mentioned	—	0.044
WOT	Most Mentioned	—	0.660
WOT	ML Classical Features	ASOIAF	0.745
WOT	ML Classical Features	SOC	0.701
WOT	ML Word Emb. Features	ASOIAF	0.699
WOT	ML Word Emb. Features	SOC	0.551

Table 2: The results of the character classifier systems.

4 Results and Discussion

4.1 Main results

The results of all the methods on both datasets are shown in Table 2. This includes the two baseline methods, and the machine learning methods with the different feature sets. We evaluate the machine learning methods using each dataset in turn as the test set and the training set.

The First Mentioned baseline is very weak. The Most Mentioned baseline is much stronger. All the machine learning methods outperform both baselines; with the exception of the SOC trained word embeddings for WOT. We attribute the poor performance of the SOC trained word embeddings, on both the WOT and ASOIAF test sets, to the small amount of training data. With only 91 chapters to generate its training cases from, and 600 dimensions of input to fit it can very easily over-fit – thus causing these bad results.

Using more training data for the ML methods generally gives better results. The datasets are, as was presented in Table 1, of very different sizes. In all but one case, the larger training dataset performed at least as well as the smaller. In the exceptional case, of the Word Embedding features testing on SOC, we note that the difference between the results for the ASOIAF and WOT training set, is getting one additional test case correct, and so may not be significant. The models deployed to our demonstration system (Section 5) use the com-

plete set of three datasets for training to achieve the best possible results.

Between the methods, the results on the ASOIAF and SOC are very strong. The results for WOT are much weaker, though still accurate enough to be useful when combined with manual checking. We believe the reason for this is that the WOT is practically a more difficult text. It comes closer to the line between third-person *omniscient*, than the more clear third-person *limited* POV of the other texts. It also has a much larger set of named entities mentioned in each chapter. As a more difficult task, features which may be clearly indicative of the main character in other texts are less indicative on WOT.

The Word embedding feature set only contains context information, not frequency information, but still performs well. Further, it only uses a single word window to each side of the named entity. This suggests that the point of view character can largely be inferred only from the usage of its name. This is in-line with a human’s ability to read an arbitrary paragraph and identify the point of view character.

Overall the performance between word embeddings and classical features is quite close. Particularly when the small SOC training set is excluded. It may be that with more training data, the word embeddings could consistently out-perform the classical features. With the limited training data, that we currently have, it seems that the classical feature set is marginally more useful. A pragmatic issue with the word embeddings is the requirement to store a large lookup table in memory (over 10GB for the FastText embeddings); though there are methods to handle this. In our prototype we allow the user to select from either of the machine learning systems or the baseline systems.

5 Demonstration System

We have deployed an online prototype to <https://white.ucc.asn.au/tools/np>. A video of its use can be found at <https://youtu.be/iu41pUF4wTY>. This web-app, made using the CherryPy framework,² allows the user to apply any of the model discussed to their own novels.

The web-app functions as was shown in Figure 1. The user uploads an ebook, and selects one of the character detection systems that we have discussed above. They are then presented with a

²<http://cherrypy.org/>

page displaying a list of sections, with the predicted main character(s) paired with an excerpt from the beginning of the section. The user can adjust to show the top-k most-likely characters on this screen, to allow for additional recall. To avoid the user having to wait while the whole book is processed, this list is dynamically loaded as it is computed. We find that the majority of the time is spent on running the preprocessing to annotate the data before the classification.

The user can select sections to retain using a series of check-boxes. They can also use a regular expression to match the character names(s) they are interested in. The sections with matching predictions character-names will be selected. As none of the models is perfect, some mistakes are likely. The user can manually correct the selection using the check-boxes before downloading the book.

6 Conclusion

We have presented a tool to allow consumers to re-structure their ebooks around the characters they find most interesting. The system must discover the named entities that are present in each section of the book, and then classify each section as to from which characters point of view the section is narrated from. For named entity detection we make use of standard tools. However, the classification is non-trivial. In this design we implemented several systems. Simply selecting the most commonly named character proved successful as a baseline approach. It is outperformed by our machine learning based approaches. While none of the classifiers are perfect, they achieve high enough accuracy to be useful.

The results are presented to the user via a web-interface. The user can use the results to select the chapters from the point of view of the character/s they are most interested in; and correct any errors the system has made. The user can then download the selected subset of their book.

A primary issue is the limited amount of labelled training data. A future version of our application will allow the users to submit corrections. This would be achieved by allowing them to select the true main character from amongst the named entities, then sorting those named entities alongside their feature vector, for use in future training.

Further tools along these lines have potential as writing aids for authors. To allow them to assess how much “screentime” is being given to each

character of their work in progress novels. With additional discourse processing, it would be possible to display other useful analytic, such as how often characters occur in the same scenes.

Another application of related work would be the automatic indexing of texts, by adding features such as the main character name, the perspective, etc. This could be applied to resources such as WikiSource, or Project Gutenberg.

Consumer fiction processing is an interesting area to develop useful tools that can change how we consume written media.

References

- Bird, S. and Loper, E. (2004). Nltk: the natural language toolkit. In *Proceedings of the ACL 2004 on Interactive poster and demonstration sessions*, page 31. Association for Computational Linguistics.
- Bojanowski, P., Grave, E., Joulin, A., and Mikolov, T. (2016). Enriching word vectors with subword information. *arXiv preprint arXiv:1607.04606*.
- Booth, W. C. (2010). *The rhetoric of fiction*. University of Chicago Press.
- Imani, M. B., Chandra, S., Ma, S., Khan, L., and Thuraishingham, B. (2017). Focus location extraction from political news reports with bias correction. In *2017 IEEE International Conference on Big Data (Big Data)*, pages 1956–1964.
- Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., and Dean, J. (2013). Distributed representations of words and phrases and their compositionality. In *Advances in Neural Information Processing Systems*, pages 3111–3119.
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., and Duchesnay, E. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830.
- White, L., Togneri, R., Liu, W., and Bennamoun, M. (2015). How well sentence embeddings capture meaning. In *Proceedings of the 20th Australasian Document Computing Symposium, ADCS '15*, pages 9:1–9:8. ACM.