

Generating Bags of Words from the Sums of their Word Embeddings

A greedy algorithms for (re-)creating the unordered
collection of words from a sum of word embeddings
representation

Lyndon White,
Roberto Togneri, Wei Liu, Mohammed Bennamoun

School of Electrical, Electronic and Computer Engineering
The University of Western Australia



What are sentence vector representations?

Methods for representing key information about a sentence, as a vector


Classical (Non-compositional): LSA, LDA, BOW ...

Compositional: RAE, RvNN, ...

Noncompositional: PV-DM, PV-DBOW, **SOWE**

We have turned sentences into numeric vectors, now we want to turn them back.

Input Sentences Manipulate Numbers Output Sentences



Why are we converting SOWE to BOW?

- ▶ Part-way step towards sentence generation.

Why are we converting SOWE to BOW?

- ▶ Part-way step towards sentence generation.
- ▶ Translating various media to keywords via common vector space.

Why are we converting SOWE to BOW?

- ▶ Part-way step towards sentence generation.
- ▶ Translating various media to keywords via common vector space.
- ▶ Theoretical implications on what information is maintained by the SOWE.

We have turned sentences into numeric vectors, now we want to turn them back.

Sentence: It was the best of times, it was the worst of times

We have turned sentences into numeric vectors, now we want to turn them back.

Sentence: It was the best of times, it was the worst of times
Vector representation: $[0.79, 1.27, 0.28, \dots, 1.29]$

We have turned sentences into numeric vectors, now we want to turn them back.

Sentence: It was the best of times, it was the worst of times

Vector representation: $[0.79, 1.27, 0.28, \dots, 1.29]$

BOW output: {best: 1, times: 2, worst: 1,
it: 2, of: 2, the: 2, was: 2, , : 1}

We have turned sentences into numeric vectors, now we want to turn them back.

Sentence: It was the best of times, it was the worst of times

Vector representation: $[0.79, 1.27, 0.28, \dots, 1.29]$

BOW output: {best: 1, times: 2, worst: 1,
it: 2, of: 2, the: 2, was: 2, : 1}

Sentence: It was the worse of times, it was the best of times

Existing methods do not produce closely matching sentences.

- ▶ Iyyer et al's compositional method
- ▶ Bowman et al's RNN based method

M. Iyyer, J. Boyd-Graber, and H. D. III, "Generating sentences from semantic vector space representations," in *NIPS Workshop on Learning Semantics*, 2014.

S. R. Bowman, L. Vilnis, O. Vinyals, A. M. Dai, R. Jozefowicz, and S. Bengio, "Generating sentences from a continuous space," *ArXiv preprint arXiv:1511.06349*, 2015.

Existing methods do not produce closely matching sentences.

- ▶ Iyyer et al's compositional method
- ▶ Bowman et al's RNN based method
- ▶ Both are demonstrated to produce loosely similar sentences.

M. Iyyer, J. Boyd-Graber, and H. D. III, "Generating sentences from semantic vector space representations," in *NIPS Workshop on Learning Semantics*, 2014.

S. R. Bowman, L. Vilnis, O. Vinyals, A. M. Dai, R. Jozefowicz, and S. Bengio, "Generating sentences from a continuous space," *ArXiv preprint arXiv:1511.06349*, 2015.

Existing methods do not produce closely matching sentences.


- ▶ Iyyer et al's compositional method
- ▶ Bowman et al's RNN based method
- ▶ Both are demonstrated to produce loosely similar sentences.
- ▶ Neither has show a demonstration on any large scale corpus.

M. Iyyer, J. Boyd-Graber, and H. D. III, "Generating sentences from semantic vector space representations," in *NIPS Workshop on Learning Semantics*, 2014.

S. R. Bowman, L. Vilnis, O. Vinyals, A. M. Dai, R. Jozefowicz, and S. Bengio, "Generating sentences from a continuous space," *ArXiv preprint arXiv:1511.06349*, 2015.

We solve the objective function to get a bag of words.

Find the inclusion vector $\tilde{c} = [c_1, c_2, \dots, c_{|\mathcal{V}|}] \in \mathbb{N}_0^{|\mathcal{V}|}$ that for we have $\min d(\tilde{s}, \sum_{\tilde{x}_j \in \mathcal{V}} \tilde{x}_j c_j)$



We solve the objective function to get a bag of words.

Find the inclusion vector $\tilde{c} = [c_1, c_2, \dots, c_{|\mathcal{V}|}] \in \mathbb{N}_0^{|\mathcal{V}|}$ that for we have $\min d(\tilde{s}, \sum_{\tilde{x}_j \in \mathcal{V}} \tilde{x}_j c_j)$

We solve the objective function to get a bag of words.

Find the inclusion vector $\tilde{c} = [c_1, c_2, \dots, c_{|\mathcal{V}|}] \in \mathbb{N}_0^{|\mathcal{V}|}$ that for we have $\min d(\tilde{s}, \sum_{\tilde{x}_j \in \mathcal{V}} \tilde{x}_j c_j)$

Input Vector $\tilde{s} = [0.79, 1.27, 0.28, \dots, 1.29]$

We solve the objective function to get a bag of words.

Find the inclusion vector $\tilde{c} = [c_1, c_2, \dots, c_{|\mathcal{V}|}] \in \mathbb{N}_0^{|\mathcal{V}|}$ that for we have $\min d(\tilde{s}, \sum_{\tilde{x}_j \in \mathcal{V}} \tilde{x}_j c_j)$

Input Vector $\tilde{s} = [0.79, 1.27, 0.28, \dots, 1.29]$

Vector Selection

$$\begin{aligned} \tilde{s} \approx & 1 \times [0.19, 0.50, 0.14, \dots, 0.59] \\ & + 2 \times [-0.15, 0.19, 0.03, \dots, -0.17] \\ & + \dots \\ & + 0 \times [0.19, 2.10, 1.34, \dots, 1.20] \\ & + 1 \times [0.79, 1.27, 0.28, \dots, 1.29] \end{aligned}$$

We solve the objective function to get a bag of words.

Find the inclusion vector $\tilde{c} = [c_1, c_2, \dots, c_{|\mathcal{V}|}] \in \mathbb{N}_0^{|\mathcal{V}|}$ that for we have $\min d(\tilde{s}, \sum_{\tilde{x}_j \in \mathcal{V}} \tilde{x}_j c_j)$

Input Vector $\tilde{s} = [0.79, 1.27, 0.28, \dots, 1.29]$

Vector Selection

$$\begin{aligned} \tilde{s} \approx & 1 \times [0.19, 0.50, 0.14, \dots, 0.59] \\ & + 2 \times [-0.15, 0.19, 0.03, \dots, -0.17] \\ & + \dots \\ & + 0 \times [0.19, 2.10, 1.34, \dots, 1.20] \\ & + 1 \times [0.79, 1.27, 0.28, \dots, 1.29] \end{aligned}$$

BOW {best: 1, times: 2, worst: 1,
it: 2, of: 2, the: 2, was: 2, : 1}

How to solve the objective function? Greedy

Find the inclusion vector $\tilde{c} = [c_1, c_2, \dots, c_{|\mathcal{V}|}] \in \mathbb{N}_0^{|\mathcal{V}|}$ that for we have $\min d(\tilde{s}, \sum_{\tilde{x}_j \in \mathcal{V}} \tilde{x}_j c_j)$

- ▶ Similarities to Knapsack family of problems.
 - ▶ Provably NP-Hard

How to solve the objective function?

Greedy

Find the inclusion vector $\tilde{c} = [c_1, c_2, \dots, c_{|\mathcal{V}|}] \in \mathbb{N}_0^{|\mathcal{V}|}$ that for we have $\min d(\tilde{s}, \sum_{\tilde{x}_j \in \mathcal{V}} \tilde{x}_j c_j)$

- ▶ Similarities to Knapsack family of problems.
 - ▶ Provably NP-Hard
- ▶ Very high dimensionality of inclusion vector
 - ▶ $|\mathcal{V}| \approx 40,000$ for Brown Corpus
 - ▶ $|\mathcal{V}| \approx 180,000$ for Books Corpus

How to solve the objective function?

Greedy

Find the inclusion vector $\tilde{c} = [c_1, c_2, \dots, c_{|\mathcal{V}|}] \in \mathbb{N}_0^{|\mathcal{V}|}$ that for we have $\min d(\tilde{s}, \sum_{\tilde{x}_j \in \mathcal{V}} \tilde{x}_j c_j)$

- ▶ Similarities to Knapsack family of problems.
 - ▶ Provably NP-Hard
- ▶ Very high dimensionality of inclusion vector
 - ▶ $|\mathcal{V}| \approx 40,000$ for Brown Corpus
 - ▶ $|\mathcal{V}| \approx 180,000$ for Books Corpus
- ▶ A greedy algorithm is linear time in $|\mathcal{V}|$

A more direct bag notation for vector selection problem.

Rather than writing:

Find the inclusion vector $\tilde{c} = [c_1, c_2, \dots, c_{|\mathcal{V}|}] \in \mathbb{N}_0^{|\mathcal{V}|}$ that for we have $\min d(\tilde{s}, \sum_{\tilde{x}_j \in \mathcal{V}} \tilde{x}_j c_j)$

A more direct bag notation for vector selection problem.

Rather than writing:

Find the inclusion vector $\tilde{c} = [c_1, c_2, \dots, c_{|\mathcal{V}|}] \in \mathbb{N}_0^{|\mathcal{V}|}$ that for we have $\min d(\tilde{s}, \sum_{\tilde{x}_j \in \mathcal{V}} \tilde{x}_j c_j)$

We can equivalently say:

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

Greedy Addition: where you add the best vector to your current bag, and repeat.

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

1. For each vector \tilde{x}_j in the vocabulary consider $d(\tilde{s}, \Sigma(\mathcal{B}) + \tilde{x}_j)$
2. Add in the vector that gets the total closest to \tilde{s}
 $\mathcal{B} \leftarrow \mathcal{B} \cup \{\tilde{x}_\star\}$
 - unless adding nothing would be better – then terminate
3. Repeat

A 1 dimensional example of greedy additon

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

Consider $\mathcal{V} = \{24, 25, 100\}$

$$\tilde{s} = 148 \quad d(x, y) = |x - y|$$

1. $\mathcal{B} = []$

$$d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 0| = 148$$

A 1 dimensional example of greedy additon

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

Consider $\mathcal{V} = \{24, 25, 100\}$

$$\tilde{s} = 148 \quad d(x, y) = |x - y|$$

1. $\mathcal{B} = []$

$$d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 0| = 148$$

2. $\mathcal{B} = [100]$

$$d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 100| = 48$$

A 1 dimensional example of greedy additon

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

- Consider $\mathcal{V} = \{24, 25, 100\}$ $\tilde{s} = 148$ $d(x, y) = |x - y|$
1. $\mathcal{B} = []$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 0| = 148$
 2. $\mathcal{B} = [100]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 100| = 48$
 3. $\mathcal{B} = [100, 25]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25)| = 23$

A 1 dimensional example of greedy additon

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

Consider $\mathcal{V} = \{24, 25, 100\}$ $\tilde{s} = 148$ $d(x, y) = |x - y|$

1. $\mathcal{B} = []$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 0| = 148$

2. $\mathcal{B} = [100]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 100| = 48$

3. $\mathcal{B} = [100, 25]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25)| = 23$

4. $\mathcal{B} = [100, 25, 24]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25 + 24)| = 1$

A 1 dimensional example of greedy additon

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

Consider $\mathcal{V} = \{24, 25, 100\}$ $\tilde{s} = 148$ $d(x, y) = |x - y|$

1. $\mathcal{B} = []$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 0| = 148$
2. $\mathcal{B} = [100]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 100| = 48$
3. $\mathcal{B} = [100, 25]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25)| = 23$
4. $\mathcal{B} = [100, 25, 24]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25 + 24)| = 1$
5. $\mathcal{B} = [100, 25, 24]$ No improvement possible $d(\tilde{s}, \Sigma(\mathcal{B})) = 1$

A 1 dimensional example of greedy additon

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

- Consider $\mathcal{V} = \{24, 25, 100\}$ $\tilde{s} = 148$ $d(x, y) = |x - y|$
1. $\mathcal{B} = []$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 0| = 148$
 2. $\mathcal{B} = [100]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - 100| = 48$
 3. $\mathcal{B} = [100, 25]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25)| = 23$
 4. $\mathcal{B} = [100, 25, 24]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25 + 24)| = 1$
 5. $\mathcal{B} = [100, 25, 24]$ No improvement possible $d(\tilde{s}, \Sigma(\mathcal{B})) = 1$

Fell for greedy trap

1-Substitution: Lessen the greed by reconsidering past choices

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

1. Consider each word vector in the current bag $\tilde{x}_a \in \mathcal{B}$
2. Would deleting it improve the score?
 $d(\tilde{s}, \Sigma(\mathcal{B}) - \tilde{x}_a) < d(\tilde{s}, \Sigma(\mathcal{B}))$?
3. Can it be swapped for another word to improve the score? $\exists \tilde{x}_b \in \mathcal{V}$ such that
 $d(\tilde{s}, \Sigma(\mathcal{B}) - \tilde{x}_a + \tilde{x}_b) < d(\tilde{s}, \Sigma(\mathcal{B}))$?

A 1 dimensional example of 1-substitution

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

Consider $\mathcal{V} = \{24, 25, 100\}$ $\tilde{s} = 148$ $d(x, y) = |x - y|$

1. $\mathcal{B} = [100, 25, 24]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25 + 24)| = 1$

A 1 dimensional example of 1-substitution

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

Consider $\mathcal{V} = \{24, 25, 100\}$ $\tilde{s} = 148$ $d(x, y) = |x - y|$

1. $\mathcal{B} = [100, 25, 24]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25 + 24)| = 1$
2. Consider swapping 100

A 1 dimensional example of 1-substitution

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

Consider $\mathcal{V} = \{24, 25, 100\}$ $\tilde{s} = 148$ $d(x, y) = |x - y|$

1. $\mathcal{B} = [100, 25, 24]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25 + 24)| = 1$
2. Consider swapping 100
3. Consider swapping 25

A 1 dimensional example of 1-substitution

Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

Consider $\mathcal{V} = \{24, 25, 100\}$ $\tilde{s} = 148$ $d(x, y) = |x - y|$

1. $\mathcal{B} = [100, 25, 24]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25 + 24)| = 1$
2. Consider swapping 100
3. Consider swapping 25
4. $\mathcal{B} = [100, 24, 24]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 24 + 24)| = 0$

A 1 dimensional example of 1-substitution


Find the bag of vectors \mathcal{B} (a multi-subset of \mathcal{V}), such that we have $\min d(\tilde{s}, \sum_{\tilde{x}_a \in \mathcal{B}} \tilde{x}_a)$

Consider $\mathcal{V} = \{24, 25, 100\}$ $\tilde{s} = 148$ $d(x, y) = |x - y|$

1. $\mathcal{B} = [100, 25, 24]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 25 + 24)| = 1$
2. Consider swapping 100
3. Consider swapping 25
4. $\mathcal{B} = [100, 24, 24]$ $d(\tilde{s}, \Sigma(\mathcal{B})) = |148 - (100 + 24 + 24)| = 0$

Fixed, but there are deeper greedy traps, that can be constructed.

Run until convergence



Experimental Setup

Preprocess corpora to only use known words.

- ▶ For word embeddings, we use pretrained GloVe

J. Pennington, R. Socher, and C. D. Manning, "Glove: Global vectors for word representation," in *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP 2014)*, 2014, pp. 1532–1543.

Preprocess corpora to only use known words.

- ▶ For word embeddings, we use pretrained GloVe
- ▶ Restrict vector vocabulary to only words used in corpora

Preprocess corpora to only use known words.

- ▶ For word embeddings, we use pretrained GloVe
- ▶ Restrict vector vocabulary to only words used in corpora
- ▶ Preprocess corpora to remove sentences with words not found in vocabulary.

We used the Brown, and the Books Corpus
as generation targets.

W. N. Francis and H. Kucera, "Brown corpus manual," *Brown University*, 1979.

Y. Zhu, R. Kiros, R. Zemel, R. Salakhutdinov, R. Urtasun, A. Torralba, and S. Fidler,
"Aligning books and movies: Towards story-like visual explanations by watching
movies and reading books," in *ArXiv preprint arXiv:1506.06724*, 2015.

We used the Brown, and the Books Corpus as generation targets.

Brown Corpus

- ▶ Extracts from 500 varied works from 1961
- ▶ 40,485 unique words
- ▶ 42,004 sentences
- ▶ Sentence Length Q3:
25 words

W. N. Francis and H. Kucera, "Brown corpus manual," *Brown University*, 1979.

Y. Zhu, R. Kiros, R. Zemel, R. Salakhutdinov, R. Urtasun, A. Torralba, and S. Fidler, "Aligning books and movies: Towards story-like visual explanations by watching movies and reading books," in *ArXiv preprint arXiv:1506.06724*, 2015.

We used the Brown, and the Books Corpus as generation targets.

Brown Corpus

- ▶ Extracts from 500 varied works from 1961
- ▶ 40,485 unique words
- ▶ 42,004 sentences
- ▶ Sentence Length Q3:
25 words

Books Corpus

- ▶ 11,038 unpublished novels
we use a random subset
- ▶ 178,694 unique words
- ▶ 66,464 sentences
- ▶ Sentence Length Q3:
17 words

W. N. Francis and H. Kucera, "Brown corpus manual," *Brown University*, 1979.

Y. Zhu, R. Kiros, R. Zemel, R. Salakhutdinov, R. Urtasun, A. Torralba, and S. Fidler, "Aligning books and movies: Towards story-like visual explanations by watching movies and reading books," in *ArXiv preprint arXiv:1506.06724*, 2015.

Results

A pair of short example

Sentence: we looked out at the setting sun .

Target BOW: . at looked out setting sun the we

Output BOW: . at looked out setting sun the we

Bowman et al's: they were laughing at the same time .

A pair of short example

Sentence: i went to the kitchen .

Target BOW: . i kitchen the to went

Output BOW: . i kitchen the to went

Bowman et al's: i went to the kitchen .

A short example where the method fails

Sentence: how are you doing ?

Target BOW: ? are doing how you

Output BOW: ? 're ~~are~~ do ~~doing~~ how well ~~you~~

Bowman et al's: what are you doing ?

A medium length example

Sentence: this is the basis of a comedy of manners first performed in 1892

Target BOW: 1892 a basis comedy first in is manners of of performed the this

Output BOW: 1892 a basis comedy first in is manners of of performed the this

Iyer et al's another is the subject of this trilogy of romance most performed in 1874

A long example

Sentence: thus she leaves her husband and child for
aleksei vronsky but all ends sadly when she
leaps in front of a train

Target BOW: a aleksei all and but child ends for front her
husband in leaps leaves of sadly she she thus
train vronsky when

Output BOW: a aleksei all and but child ends for front her
husband in leaps leaves of sadly she she thus
train vronsky when

Iyer et al's: however she leaves her sister and daughter
from former fiancé and she ends unfortunately
when narrator drives into life of a house

A long example where the method fails.

Sentence: ralph waldo emerson dismissed this poet as the
jingle man and james russell lowell called him
three-fifths genius and two-fifths sheer fudge

Target BOW: and and as called dismissed emerson fudge
genius him james jingle lowell man poet ralph
russell sheer the this three-fifths two-fifths waldo

Output BOW: **2008** _..._(13) _..._(34) _..._(44) “ **aldrick** and
and ~~as~~ **both** called dismissed emerson fudge
genius **hapless** him **hirsute** james jingle **known**
lowell man poet ralph russell sheer the this
three-fifths two-fifths waldo **was**

A long example where the method fails.

Sentence: ralph waldo emerson dismissed this poet as the
jingle man and james russell lowell called him
three-fifths genius and two-fifths sheer fudge

Target BOW: and and as called dismissed emerson fudge
genius him james jingle lowell man poet ralph
russell sheer the this three-fifths two-fifths waldo

Iyer et al's: henry_david_thoreau rejected this author like
the tsar boat and imbalance created known
good writing and his own death

Yet another example

Sentence: in a third novel a sailor abandons the patna
and meets marlow who in another novel meets
kurtz in the congo

Target BOW: a a abandons and another congo in in in
kurtz marlow meets meets novel novel patna
sailor the the third who

Output BOW: a a abandons and another congo in in in
kurtz marlow meets meets novel novel patna
sailor the the third who

Iyer et al's: during the short book the lady seduces the
family and meets cousin he in a novel dies sister
from the mr.

A final example

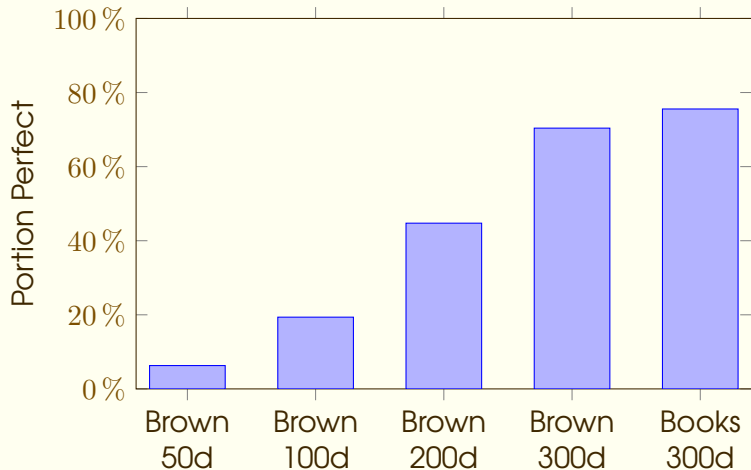
Sentence: name this 1922 novel about leopold bloom
written by james joyce

Target BOW: 1922 about bloom by james joyce leopold
name novel this written

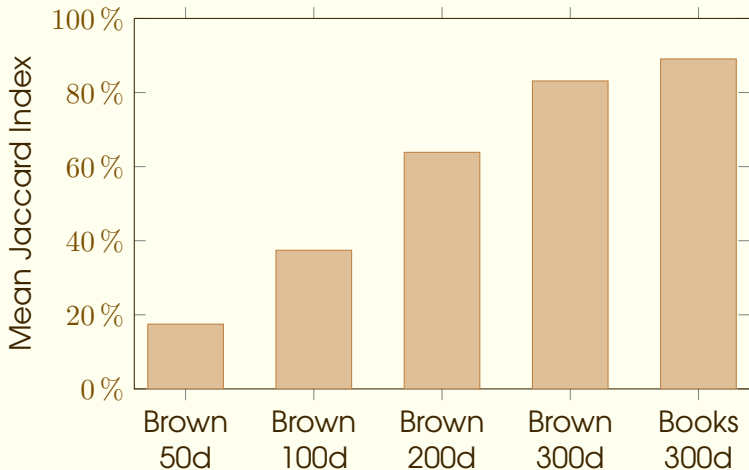
Output BOW: 1922 about bloom by james joyce leopold
name novel this written

Iyer et al's: name this 1906 novel about gottlieb_fecknoe
inspired by james_joyce


Results: The more dimensions used in the word embeddings, the better the recovery.



Results: The more dimensions used in the word embeddings, the better the recovery.




Result: The longer the sentence, the worse recovery



Brown Q3: 25 words

Books Q3: 17 words

Result: The larger the vocabulary, the worse recovery



Brown $|\mathcal{V}| \approx 40,000$

Books $|\mathcal{V}| \approx 180,000$

Conclusion

Future Work: we could to order them to get a sentence.

- Use a language model to find probability of any given sequence.


Future Work: we could to order them to get a sentence.

- ▶ Use a language model to find probability of any given sequence.
- ▶ Not guaranteed to find a single unique order.


Future Work: we could to order them to get a sentence.

- ▶ Use a language model to find probability of any given sequence.
- ▶ Not guaranteed to find a single unique order.
- ▶ Also NP-hard.

A two step method for generating sentences.



A two step method for generating sentences.



Conclusion: We can often successfully recover the BOW, from the SOWE

- ▶ Vector selection with a greedy algorithm
 - ▶ This is a broad generalisation of Knapsack Problem
 - ▶ Input: SOWE vector
 - ▶ Greedy Addition + 1-Substitution til convergence.
 - ▶ Output: BOW

- ▶ Future work: order the words using a language model.

Appendix

Recent results suggest sum of word embeddings captures surprising amounts of semantic information

Category	Example
Adhesion to Vertical Surface	There is a magnet on the refrigerator.
Support by Horizontal Surface	There is an apple on the refrigerator.
Support from Above	There is an apple on the branch.
Full Containment	There is an apple in the refrigerator.
Partial Containment	There is an apple in the water.

- ▶ Categorise sentences based on the positional component of their meaning.
- ▶ Ritter et. al. found sum of word embeddings to outperform all more complex models.

S. Ritter, C. Long, D. Paperno, M. Baroni, M. Botvinick, and A. Goldberg, "Leveraging preposition ambiguity to assess compositional distributional models of semantics," *The Fourth Joint Conference on Lexical and Computational Semantics*, 2015.

Recent results suggest sum of word embeddings captures surprising amounts of semantic information

- ▶ We groups MSRP and Opinions sentences by semantic equivalence forming classes of paraphrases.
- ▶ Then used various sentence embeddings as input to a linear SVM to try and classify back into the groups.
- ▶ SOWE was amongst top contenders (<0.6% worse than best in both cases)

Iyyer et al's compositional sentence generation method.

- ▶ Variation on the URAE
- ▶ Reuses a neural network to (merge up the dependency tree
- ▶ Similar to unfold.
- ▶ Requires structure of output to be given as a input.

Bowman et al's RNN based sentence generation method.

- ▶ Use LSTM RNN for decode/encoding step
- ▶ Use VAE as representation of posterior probabilities.
- ▶ lots of interesting properties and other uses.

Results

Corpus	Word Embedding Dimensions	Portion Perfect	Mean Jac-card Score	Mean Precision	Mean Recall	Mean F1 Score
Brown	50	6.3%	0.175	0.242	0.274	0.265
Brown	100	19.4%	0.374	0.440	0.530	0.477
Brown	200	44.7%	0.639	0.695	0.753	0.720
Brown	300	70.4%	0.831	0.864	0.891	0.876
Books	300	75.6%	0.891	0.912	0.937	0.923